

MANUAL

Every knob in AlphaNode, explained in plain language.

What each panel shows, what every setting actually changes, and what each leaderboard column measures — with the shipped default, the range the app allows, and the trap worth knowing before you turn it up. One entry per knob, numbered, so you can jump straight to the one in front of you.

CONTENTS

01 The basics — what you mine, and what the main window shows

1.01	A formula (an "alpha")	3
1.02	A round	3
1.03	TRAIN / VAL / TEST	3
1.04	Fitness — the number the search maximises	4
1.05	► Start node	4
1.06	■ Stop	4
1.07	● stopped / starting / running (the state pill)	5
1.08	⚙ Settings	5
1.09	📅 Sessions	5
1.10	Light / Dark	5
1.11	Activate node / Unlock	6
1.12	ROUNDS · FORMULAS TRIED · ALPHAS FOUND	6
1.13	BEST FITNESS	6
1.14	The live log	6
1.15	SIGNAL API — running services	7
1.16	LEADERBOARD	7
1.17	PORTFOLIO — top-N combined	7
1.18	FORWARD TRACK — paper steps	8

02 Pairs, timeframe and market data

2.01	Which pairs to trade	9
2.02	Timeframe (bar size)	9
2.03	Market data (the automatic download at Start)	10
2.04	Data and state on disk	10

03 The search settings — how hard and how long it looks

3.01	Resources (CPU share)	11
3.02	Population	11
3.03	Generations	11
3.04	Seed (0 = auto)	12
3.05	Pause, sec	12
3.06	Status port	12
3.07	Optimize by win rate	12
3.08	Explore every N-th	13
3.09	Max. rounds (0 = ∞)	13
3.10	Leaderboard size	13
3.11	Warm-start from library	14

04 The engine knobs — simulation, formula shape, breeding, fitness

4.01	Target-vol (ann.)	15
4.02	Fee (turnover)	15
4.03	Max. depth	15
4.04	Max. nodes	16
4.05	Tournament size	16
4.06	Elitism	16

4.07	Random/generation	16
4.08	Crossover share	17
4.09	Complexity penalty	17
4.10	Correlation threshold	17
4.11	Similarity penalty	18
4.12	Hall of Fame size	18
4.13	Robust blocks (0 = legacy)	18

05 Date segments, and the honesty discipline

5.01	TRAIN start	19
5.02	VAL start	19
5.03	TEST start	20
5.04	TEST end	20
5.05	The recommended cut at each bar size	20
5.06	Working with a re-cut	21
5.07	by — how portfolio members are picked	21

06 The leaderboard, column by column

6.01	★	22
6.02	#	22
6.03	fitness	22
6.04	TEST OOS	22
6.05	maxDD	22
6.06	CAGR	23
6.07	sortino	23
6.08	T ↑	23
6.09	T ↓	23
6.10	T ~	23
6.11	L/S /yr·a	24
6.12	tr/yr·a	24
6.13	win%	24
6.14	win ↑	24
6.15	win ↓	24
6.16	ID	25
6.17	formula	25

07 What you can do with a formula

7.01	Double-click a row → the equity window	26
7.02	Passport	26
7.03	PDF report	26
7.04	Serve signal (API)	26
7.05	Download signals (CSV)	27
7.06	Forward track +	27
7.07	Chart (Forward Track panel)	27
7.08	★ Favorites	27
7.09	Copy formula / Copy formula + metrics	28
7.10	Right-click a header → the Columns menu	28
7.11	Sorting	28
7.12	Export table (CSV)...	28
7.13	Export full library (CSV)...	29
7.14	► Build portfolio	29
7.15	"families only" switch	29
7.16	Clear all history	29

Part 01

The basics — what you mine, and what the main window shows

What an alpha is, what one round does, how history is split, and what every card on the main window is telling you. Nothing here is a knob you have to tune.

1.01 A formula (an "alpha")

A short maths expression over the price tables — a real one looks like `ema:20(ts_delta:14(cs_zscore(cs_demean(add(ts_max:30(volume),ret))))`). The number after a colon is a lookback in bars. It is evaluated once per coin per bar: positive means be long this coin, negative short, bigger magnitude a bigger share of the book. Leaves are the only 12 things it can see (close, open, high, low, volume, ret, vwap, range, body, dvol, logret, funding); branches are 25 operators — arithmetic, shape-benders (`neg/sign/abs/slog/tanh`), lookbacks over 2–200 bars of one coin's own history (`ts_mean`, `ts_std`, `ts_zscore`, `ts_min`, `ts_max`, `ts_delta`, `ts_delay`, `ts_sum`, `ts_roc`, `ema`) and cross-sectional ones that compare all your coins on the same bar (`cs_rank`, `cs_zscore`, `cs_demean`, `cs_scale`). Everything downstream is fixed and identical for every alpha: divide by the coin's recent volatility, normalise to 100% of the book, scale to your target vol, hold through a 10% no-trade band, pay fees on turnover. So the search can only change the SIGN and the RELATIVE SIZE of the numbers across coins and across time.

Watch out. Absolute scale is a no-op — `add(x,x)` is exactly the strategy `x`, two nodes more expensive. And cross-sectional operators need a real cross-section: with 1–2 pairs, `cs_rank` and `cs_demean` collapse to constants.

1.02 A round

default **population 200 · 25 generations · Hall of Fame 15 · pause 5 s · explore every 4th round**

range **population 4–4,000 · generations 1–500 · HoF 1–100 · pause 0–3,600 s**

One complete evolutionary search, start to finish: build 200 candidate formulas, score them all, breed the survivors for 25 generations, fine-tune the lookback windows of the top 5 finishers, then hand back up to 15 champions kept deliberately un-alike. New ones are appended to your library; the node pauses and starts again with a new seed. A round costs at most population × generations evaluations and usually fewer — repeats inside a round are scored once. Rounds are the unit of progress: let it run for tens of them before judging anything.

Watch out. Warm-start only ever sees champions from the CURRENT session, and only the top few ("Leaderboard size", 20) of them. Before you activate the node everything on disk is sealed, so the refine pool starts empty and refine may never run at all.

1.03 TRAIN / VAL / TEST

default **TRAIN 2019-09-05 · VAL 2021-11-01 · TEST 2023-01-01 · TEST end 2026-07-05**

History cut into three consecutive ranges. TRAIN is where evolution works. VAL is a second era that breeding never optimises directly — it catches formulas that merely memorised TRAIN. TEST is held out entirely: not the fitness, not the selection, not the seeding ever reads it; its numbers are carried along

purely so you can look afterwards. Treat TEST as a report card you read once you have decided what to keep — the moment you pick alphas BECAUSE their TEST is high, it stops being out-of-sample, which is what every $\triangle$ on a TEST-flavoured action means.

Watch out. TEST end is a fixed calendar date. Everything after it falls outside all three segments and is never simulated — so if you never edit that field, the fresh data Start keeps downloading does nothing.

1.04 Fitness — the number the search maximises

```
default min(TRAIN, VAL) Sharpe · complexity 0.010/node · correlation 0.70 · similarity 0.5 ·
concentration 0.3 under 3 positions · activity floor 10%
```

The WORSE of two Sharpe ratios: $\min(\text{TRAIN Sharpe}, \text{VAL Sharpe})$ — a formula is only as good as its weakest era, so curve-fitting one bull market buys nothing. While breeding, the search also subtracts formula size, correlation above 0.70 with a champion already found, and concentration when the book averages fewer than 3 meaningful positions. A formula in the market on under 10% of TRAIN or of VAL bars is discarded unscored. Fitness is the only score not contaminated by TEST — but a high one is no promise: rows scoring above +2 in the shipped library go on to a negative held-out TEST.

Watch out. The fitness you SEE is not the one the search ranks by: the stored number carries only the concentration penalty, while size and diversity steer breeding and are then thrown away. That is why the log's green $\star$ line reads lower than the round summary for the very same formula.

1.05 ► Start node

Saves your settings, then checks your market-data snapshot: if it is missing, missing one of your pairs, or its newest bar is over 5 days old, a console window opens, downloads exactly your pairs at your timeframe from Binance, and Start resumes by itself. Then it launches the node as a separate process, hands it every setting as an environment variable and polls its status every 1.5 s. A ticker that is not a real Binance USDT-perp is dropped with a warning rather than looping. The node also publishes a live status page on <http://localhost:8787> — bound on all interfaces, so anyone on your LAN who knows the port can read it.

Watch out. The node reads your settings ONLY at launch. Editing population, dates, fees or the universe while it runs changes nothing until you Stop and Start again.

1.06 ■ Stop

Sends the node an interrupt. It notices at the next generation boundary, writes its final status and exits — seconds to a couple of minutes. Stop before anything that rewrites your library: activating, loading a session and clearing history all refuse to run while a miner is appending to the same files. Nothing already written is lost.

Watch out. The tooltip says the current round will finish — it does not. The round in flight is abandoned and its whole Hall of Fame discarded before reaching the library. If the log shows a $\star$ new best, wait for "✓ round N done" before stopping.

1.07 ● stopped / starting / running (the state pill)

What the node itself reports: grey stopped, blue starting (booting — checking or downloading data, loading your library), green running. Beside it, a one-line resource summary: CPU budget, cores that works out to, universe and effective target volatility. When nothing runs, those lines are prefixed "last run —" so you cannot mistake stale text for telemetry. "starting" can legitimately last a while on the first run of a timeframe.

Watch out. The pill mirrors the node's status FILE, not the process. After a hard kill it keeps saying green "running" with nothing behind it — the giveaway is that the Start button is clickable.

1.08 ⚙ Settings default hidden on a fresh install

Shows and hides the left-hand panel holding every knob the engine understands. Hidden on a fresh install, remembered between launches, always exactly as wide as its content (deliberately not draggable). Every field carries a hover tooltip. At the bottom sit "Reset to defaults" and the red "Clear all history". Leave it closed for normal use — the defaults are what the engine was tuned with.

Watch out. Reset keeps two things on purpose: your theme and whether this panel is open. And "Clear all history" deletes the libraries and round history of EVERY timeframe at once, plus the built portfolio — settings, favourites and the forward track survive, your mined alphas do not, and nothing is backed up first.

1.09 📁 Sessions

Your whole workspace as one .tar.gz — everything on the board. Every timeframe's library (so the leaderboard comes back as it was), the round history, the forward track, the built portfolio with its members, your ★ favourites, the run status — the four counters at the top, the live log and the round ticker — and your settings. Two things are deliberately left out: market data, which re-downloads, and your subscription key, so a session travels safely to another machine or another person. The registry of running signal services stays behind too — it is a list of process ids on this machine and means nothing anywhere else. Every save mints its own 6-character ID, the same shape the leaderboard uses for an alpha. It is in the manifest and in the filename, so an archive sitting in a backup folder or arriving by mail says which session it is without being opened — and two sessions with the same name, which is the normal case, are finally distinguishable. Loading replaces the current workspace and refuses to run while the node, a forward step or a portfolio build is working. Double-click a row to see what is inside before you load it.

Watch out. A session is saved mid-run, so its status says the node was running; loading rewrites that one field to "stopped" and dims the line to "last run —", because nothing is actually searching after a load. The counters and log it carries are history and are kept verbatim. The tooltip promises an auto snapshot every time the node stops. There is none — the only automatic snapshot in the app is the before-load checkpoint, so the one safety net you get is for the operation you were already worried about.

1.10 Light / Dark default follows the operating system on the very first launch

Repaints the palette and rebuilds the window from scratch — that rebuild is what lets the parts which cannot restyle themselves live (the table, the equity images, the chart canvas) come back in the right colours. Your leaderboard, chart and running services are restored immediately after. The flicker is by

design. Harmless to flip at any time, including mid-search.

Watch out. "Follow the OS" happens once, at startup: the resolved value is written into your settings and saved by the next Start. There is no system position to go back to, and Reset preserves the theme — the only way back is editing `gui_settings.json` by hand.

1.11 Activate node / Unlock

Every formula the node mines is sealed before it touches your disk: the library stores encrypted text plus a public 12-character id, and the row shows its stored metrics but "locked" instead of the maths. Activate takes your subscription key once, claims a machine seat, and unlocks every sealed formula in every local library in one pass; from then on the node mines in the open, re-checked at each Start. Do it early — sealed rows are inert to everything that has to re-simulate: the computed columns, the Portfolio build, Serve and the forward track all need the plain text.

Watch out. Activation rewrites the library files, so it refuses while the node is mining — stop, activate, then start. And the per-row "Unlock" is session-only: the file on disk stays sealed and the row is locked again after a restart.

1.12 ROUNDS · FORMULAS TRIED · ALPHAS FOUND

The three tiles, all read from the node's status file rather than your library. ROUNDS is how many complete searches have finished, resumed from the round history across restarts. FORMULAS TRIED is the lifetime total of distinct formulas evaluated — the honest measure of compute spent. ALPHAS FOUND is how many distinct formulas are in the library being mined into. Watch ALPHAS FOUND against ROUNDS to judge productivity: once a library is mature, whole rounds keeping nothing is normal, not a fault.

Watch out. Switching timeframe does NOT move these tiles — the table below re-reads that library at once, the tiles keep the last run's numbers until you press Start. FORMULAS TRIED is genuinely cross-timeframe: it comes from one shared status file.

1.13 BEST FITNESS

The fitness of the current number-one champion — the same `min(TRAIN, VAL)` figure, drawn in the accent colour because it is the number that says whether the search is getting anywhere. It only moves at round boundaries, so a mid-round silence is not a stall. Expect a jump in the first handful of rounds and then a creep; a long flat stretch means the easy structure is exhausted at these settings — widen the universe, raise the complexity limits, or move to the Portfolio card. Shows "—" until the first round finishes, and never shows a TEST number.

Watch out. In win-rate mode this tile still formats the value as a Sharpe-style "+0.56" instead of "56%" — the node always writes 'sharpe' into its status file. The leaderboard's fitness column tags each row from the row itself, so trust the column over the tile.

1.14 The live log

Seven colour-coded lines from the node's own event feed (the last 80 events are kept): round starts and finishes in normal ink, ★ new bests in green, window fine-tuning in blue, warnings and errors in red. "► round N: explore" versus "refine" tells you the mode; "✓ round N · 47s · 4,812 formulas tried · +3

champions kept · best fitness +1.81 · held-out TEST −0.26" gives you the cost and the yield of a round in one line. This is the best single place to see whether the node is healthy.

Watch out. While the library is sealed, formula text is stripped out of the log on purpose — the ★ line keeps its numbers but the maths after the em-dash is cut. Nothing is broken; the vault is doing its job.

1.15 SIGNAL API — running services

A card that appears once you are serving something. "Serve" starts a small local web service that recomputes that strategy's live target positions and hands them out as JSON, on the next free port from 8799; each gets a row with its URL, log file, health and a way to shut it down. This is the hand-off to your own execution code — poll the URL, get target weights per asset, trade them however you like. The services are detached: they keep serving after you close AlphaNode, and the app re-adopts them next launch by probing ports 8799–8808.

Watch out. Because a service outlives the GUI, a stale one from a previous session can still be answering on 8799 with an old strategy. Check the row's label and started-time before blaming the formula.

1.16 LEADERBOARD default every alpha at this timeframe, sorted by fitness, families-only off

The library itself as a table — every formula ever mined at this timeframe, not the node's top 20 — re-read in a background thread whenever the file changes. Click a header to sort, double-click a row for its equity chart, right-click for copy, CSV export and a column picker. The search box filters live by id or formula text. "Families only" collapses the table to one representative per formula shape, the fastest way to see how many genuinely distinct ideas you own. Rows are tinted green only when their held-out TEST Sharpe is at least zero — in a healthy library that is a minority, and that is the honest picture.

Watch out. Clicking "fitness" or "TEST OOS" does more than re-sort — it re-selects WHICH alphas are loaded from the library, and the heading changes to say so. On a sealed library only the numbers stored at mining time survive: sortino, $T\uparrow/T\downarrow/T\sim$, L/S, trades and win% show "—" for every row.

1.17 PORTFOLIO — top-N combined

default top 6, by TEST · 2–20 (a widget default: not stored in settings, untouched by Reset)

Runs N alphas from your library together and plots the combined equity with its metrics above it: Sharpe on TRAIN, VAL and TEST separately, TEST CAGR and max drawdown, against buy-and-hold. This is where mining hundreds of alphas pays off — the combined curve is usually far better than any single member, because their mistakes are uncorrelated. The "by" selector picks the members: TEST Sharpe, fitness, or "combo" (a greedy search for the best COMBINATION, optimised on TRAIN+VAL only). Daily uses the full simulator, one to two minutes; intraday uses the fast approximation and finishes in seconds — the line under the chart always names which engine produced the numbers. Under the metrics sits the member list: one row per alpha that went in, in pick order, with its ID, its SOLO TEST Sharpe (that member on its own), and the fitness and TEST OOS the leaderboard shows for the same row. Members are equal-weight — none is sized larger than another. Rows the combined Sharpe beat are tinted; that tint IS the diversification gain, and a row without it is a member the mix would be better off without. Double-click one for its own equity chart. Past twelve members the list scrolls instead of growing the card.

Watch out. The 2-20 range is enforced on the built portfolio, not on the box: the spinner accepts a typed 150, and the status line then tells you it is building 20. The range is not arbitrary — "combo" caps its search pool at 30, so asking for a combination larger than that would stop searching and hand back the whole pool. Note also that you get the number you asked for only if the library holds that many DISTINCT alphas: near-clones are filtered at 85% similarity, so a 45-row library yields about 39 members however high you set the number. "by TEST" is the dishonest pick and the card says so: the window that chose the members is the window being scored. Prefer fitness or combo. The builder also skips every sealed row, so an un-activated library stops with "need at least 2 scored alphas" however many rows you can see.

1.18 FORWARD TRACK — paper steps

default starts at \$10,000 · one step per closed bar · checked every 5 minutes

The honest verdict machine. Enrol a portfolio or a single formula and it is frozen — formulas, universe, target volatility, fee — then paper-traded forward from \$10,000, one step per closed bar: fetch the latest closed candles, recompute targets with the real engine, mark to market, rebalance, charge fees. History is append-only; nothing is recomputed backwards. This is the only number in the app that no amount of peeking can contaminate, because the bars did not exist when the formula was mined. Enrol your best candidates early and give them months. The node steps the track every 5 minutes while it mines, so it advances with the window closed.

Watch out. "Delete" is permanent and takes the paper history with it — re-enrolling the same formula starts a fresh \$10,000 track rather than resuming, so a year of forward evidence goes in one click.

Part 02

Pairs, timeframe and market data

The only settings that change the DATA rather than how it is searched. Get them right before touching anything else. AlphaNode downloads its own history from Binance USD-M perps — there is no file to point it at and no Download button.

2.01 Which pairs to trade default BTCUSDT, ETHUSDT, SOLUSDT, XRPUSDT, BNBUSDT

range no limit is enforced anywhere in the app

The exact basket the whole app runs on — search, leaderboard statistics, equity charts, CSV, live signals and the forward track use this list and nothing else. Full Binance USDT-perpetual tickers only (BTCUSDT, not BTC): what you type is uppercased and de-duplicated, but nothing appends the USDT for you. Write as many pairs as you like and press Enter once — commas, spaces, tabs and line breaks all separate them, and the whole box turns into chips. Nothing commits while you type: a comma is just a character, so the text stays exactly as you wrote it until Enter (clicking away or pressing Start commits it too, so a pair left in the box is never lost). A paste commits itself, no Enter needed; × removes a chip; clicking a chip's text pulls it back into the box to fix a typo (it returns at the END of the list); Backspace in an empty box pulls the last chip down. Below about 5–6 pairs the cross-sectional operators the search leans on stop meaning anything — that is the practical floor. Going wider buys a broader cross-section and costs download, memory and simulation time in direct proportion. A recently listed coin is safe to add: its pre-listing bars stay blank and are masked out.

Watch out. A one-pair universe looks like it works and does not — the weight is then always exactly ± 1 , so only the SIGN is traded and every alpha collapses to the same flip. And library rows do not record which basket they were mined on: change the pairs and one row mixes two worlds — fitness, TEST OOS, maxDD and CAGR are the OLD basket's stored numbers while win%, L/S, sortino, the T columns and the chart are re-simulated on the new one.

2.02 Timeframe (bar size) default 1d range 1d | 4h | 1h | 15m

The bar size for the entire pipeline. Every window inside a formula is counted in BARS of this size, so a 20-bar average is 20 days on 1d and 5 hours on 15m; Sharpe is annualised over 365 days, because crypto trades 24/7. Each timeframe is a separate world — its own snapshot file, its own alpha library, its own round history — and alphas from different bar sizes never mix or compete. Start on 1d: longest history, lightest download, fastest generations, and the only snapshot that ships with the app. Drop to 4h or 1h for more bars per year of calendar; intraday history is shorter, the download far heavier, and fees eat a much larger share of each bar's move. The choice takes effect only when settings are written — press Start, or collapse the settings pane. The line under the dropdown states that bar size's ceiling on the whole TRAIN-to-TEST-end window, which is what stops a 15m search from being pointed at six years of history.

Watch out. Picking from the dropdown OVERWRITES all four date fields with that bar size's recommended dates — including when you re-pick the timeframe you are already on. Hand-tuned dates are thrown away silently. Always pick the bar size first, then edit the dates.

2.03 Market data (the automatic download at Start)

Start inspects the snapshot for the active timeframe and downloads first if it is missing, unreadable, missing a pair from your list, or its newest bar is over 5 days old. A console window streams the downloader's live output and Start resumes by itself when it exits cleanly. If Binance's live API is unreachable from where you are, it falls back on its own to the public data archive — the same bars, 10–30 hours behind. The wait scales with bar size: 6 pairs in parallel on 1d, 4 on 4h, 3 on 1h, 2 on 15m — deliberately slow, because more parallelism trips the rate limit. Closing the console cancels safely: the file is written under a temporary name and swapped in only on complete success. History always begins at the timeframe's recommended date or the pair's listing date, whichever is later, so setting TRAIN start earlier buys no extra data. Funding-rate history comes with the candles, which is what makes the funding term mean anything.

Watch out. This is a replacement, not a top-up: the app asks for exactly your current pairs and the whole snapshot is overwritten with just those. Adding one pair re-downloads all of them from scratch, and a pair you removed loses its stored history at the next refresh.

2.04 Data and state on disk

In the installed app everything writable lives in one folder — `~/.local/share/AlphaNode` on Linux, `%APPDATA%\AlphaNode` on Windows, `~/Library/Application Support/AlphaNode` on macOS. It holds the price snapshots, `gui_settings.json`, `exports/`, and `state/` with the mined alphas (`library.jsonl`, `library_4h.jsonl` ...), the round history, `status.json`, the portfolio, the forward track, favourites, sessions and this install's node id. Your alphas are just those `library*.jsonl` text files — back them up, or use Sessions → Save current.... Deleting a `data_*.pickle` is harmless; the next Start downloads it again. Settings are written only when something triggers a save: Start, a theme switch, collapsing the settings pane, and most layout fiddling all rewrite the whole file.

Watch out. Closing the window does NOT save — a pairs or timeframe edit followed straight by quitting is simply lost. And "Clear all history" counts only the active timeframe in its dialog, then deletes every timeframe's library and history, plus `status.json`, the portfolio and the cached equity images. Nothing is backed up automatically.

Part 03

The search settings — how hard and how long it looks

How hard the app hunts, for how long, and how it splits its time between finding new ideas and polishing old ones. All of it is read once, at Start. None of it can touch the honesty of the result — TEST is held out of selection whatever you set here.

3.01 Resources (CPU share) default 50% range 5–95 (slider only)

The share of your processor handed to the search. The bold line above the slider does the arithmetic live — the share times your core count, rounded to whole worker processes, never fewer than one: "50% → 4 of 8 cores". This is the single biggest lever on wall-clock speed. It costs memory (every worker builds its own copy of the price panel, so RAM grows with the worker count), heat, and a machine that feels sluggish. Lower it to about a third to keep working on the same computer. On Linux and macOS the node also asks for background priority; on Windows that call is skipped.

Watch out. The worker pool is created fresh every round and each worker rebuilds the price panel from scratch, so a big core share buys little when Generations is very small — you pay that setup cost per round instead of spreading it over a long one.

3.02 Population default 200 range 4–4,000

How many candidate formulas exist side by side in one generation. Generation 0 is random trees (plus your warm-start champions); every later generation is rebuilt from the previous by selection and breeding. Raise it for a wider sweep of formula space — more building blocks in play, less chance of getting stuck polishing one mediocre idea. Cost is close to linear in time. If you want more total search, raise Population before Generations: breadth is what protects you from settling on the first decent idea the run stumbles into.

Watch out. Each generation is assembled as elites, then random injections, then bred children until the count is reached — nothing trims it back. Any Population below Elitism + Random/generation (16 with the defaults) is silently overridden. A number typed into the box is also not clamped; only the arrows respect the range.

3.03 Generations default 25 range 1–500

How many cycles of scoring, selection and breeding happen inside ONE round. After the last generation comes a short extra pass over the top five: each lookback window is tried at 0.8× and 1.25×, up to three passes, keeping a change only if it improves the same TRAIN/VAL score. Raise it to let a round dig deeper into what it has found; the returns flatten, and you can watch them flatten — if the ★ new-best lines dry up around generation 15, there is little point paying for 40. Lower it for more, shorter rounds: more from-scratch restarts and more variety, at the cost of paying the per-round setup more often. Below about 10, evolution barely has time to combine anything.

Watch out. evolution/config.ini says generations = 30 and the app never uses it. The node always overrides population, generations, seed and cores with the panel's numbers (or 200 / 25 headless). Only the standalone runner reads that file.

3.04 Seed (0 = auto) default 0 range 0–999,999

The starting number for the random generator that builds and mutates formulas. The same seed with the same data and settings reproduces the same run exactly. At 0 the app mints a seed once from a random node id in its state folder, so your install walks its own path through formula space and no two installs mine identical libraries. Leave it at 0 — that is what keeps your library yours. Set an explicit number only for repeatability: filing a bug, or A/B-testing one setting with the randomness held still. There is no quality difference between seeds; a "good seed" is survivorship bias.

Watch out. A shared number is a shared search — anyone running the same integer on the same settings and data mines the same formulas. Reproducibility also only holds for explore rounds: refine rounds start from whatever champions are in memory, so two machines with the same seed and different libraries diverge from the first refine onwards.

3.05 Pause, sec default 5 range 0–3,600

How long the node idles between rounds. It sleeps in half-second slices and checks for Stop before each, so Stop still responds during the pause. Raise it on a laptop or a thermally limited machine to let the fans catch up; it is pure lost mining time and has no effect on what is found. Set 0 to chain rounds back to back.

Watch out. This is not a load limiter. During a round the node still uses every core the CPU-share slider gave it, at full tilt — if the machine is too loud, lower Resources, not this.

3.06 Status port default 8787 range 1024–65535

The TCP port for the node's own small web page — a live table of the top formulas, the round log and the forward track at <http://localhost:PORT>, reloading every four seconds, with the raw numbers at [/status.json](#). The desktop window does not use this port at all (it reads the status file from disk), so the page is purely a bonus: a way to check on a mining machine from your phone or a headless server. Leave it alone unless something else owns 8787 or you want two nodes at once.

Watch out. The page listens on ALL network interfaces — anyone on your network who can reach the port can open it, and while sealed formulas show as locked, the timings, fitness and held-out TEST numbers are in the clear. If the port is taken the node keeps mining silently: no page, no dialog, nothing in the log. Keep clear of 8799 and the fifty ports above it — that is where live-signal services live.

3.07 Optimize by win rate default off

Switches what the search maximises. Unticked it maximises Sharpe — return per unit of wobble. Ticked it maximises the per-bar win rate: out of the bars where the strategy actually held a position, the share that closed in the black (flat bars count as neither). Either way it is the WORSE of TRAIN and VAL, so a formula must hold up on both. The score used for selection is deliberately lower than the raw percentage: the share is pulled toward 50% by the square root of how often the formula was active (62% of 40 bars is weaker evidence than 55% of 700), then one binomial standard error is subtracted with a small prior, so a perfect 5-of-5 reads as luck. Under 30 active bars in TRAIN or VAL, the formula leaves the round. Complexity and similarity penalties are automatically scaled down tenfold here, because win rate lives

near 0.5 while Sharpe spans a range several times wider. Use it when you want an equity curve that feels steady day to day, or as a second opinion when every Sharpe champion looks like the same trend-following family. Takes effect at the next Start.

Watch out. Win rate is blind to SIZE: 65% small green bars against 35% catastrophic red ones is a losing strategy this objective happily ranks first — judge anything it mines by TEST Sharpe and maxDD. In this build the switch also fails to reach the node's own ranking of the library, so its top list, the :8787 page and the fitness tile stay Sharpe-ordered, and refine rounds re-seed from Sharpe formulas. Mine win-rate alphas into a fresh library, or a timeframe whose library is still empty.

3.08 Explore every N-th default 4 range 1–100

The mix between hunting and polishing. Every round whose number divides evenly by N explores — a completely fresh random population, looking for families you have never seen. Every other round refines: the population starts from the champions on the node's top list. With 4 the pattern is refine, refine, refine, explore... Each round announces which it is in the live log. Lower N means a broader, more varied library but a best-fitness that creeps; higher N climbs the headline faster but converges onto one or two ideas — and a library of eight variations on one formula is one strategy, not eight. If the top of your leaderboard has been the same shape for hours, lower it; if you have hundreds of scattered mediocre alphas, raise it.

Watch out. N = 1 quietly turns refinement off completely — every round number divides by 1, so everything explores and the warm-start switch stops doing anything. The only visible symptom is that every round in the log says "explore".

3.09 Max. rounds (0 = ∞) default 0 range 0–999,999

Stops the node automatically after this many rounds; 0 means it never stops on its own. Use a finite number for a bounded experiment or an overnight budget you do not want to babysit. It is a count of rounds, not of time — the same number can mean twenty minutes or twenty hours depending on population, generations, pairs and cores.

Watch out. The count is lifetime, not per session: the node restores its round number from your history, so with 120 rounds recorded and Max rounds 50 it starts, sees it is already past, and stops without mining anything. Set it above the round count on the dashboard.

3.10 Leaderboard size default 20 range 1–200

How many champions the running node keeps in its own top list — what the :8787 page shows, what the best-fitness tile is drawn from, and, most importantly, the exact pool a refine round warm-starts from. Read it as "how wide is my warm start". Small (5–10) makes refine hammer on a handful of formulas: fast convergence, narrow results. Large (50+) gives a broad, varied starting pool. It also decides how much of an existing library is eligible at all: on start the node keeps only this many rows in memory, and everything below the cut can never seed a refine round again. Keep it comfortably below Population — an order of magnitude smaller is sane.

Watch out. It does not limit your library and never deletes anything; the app's table reads the library file, not this list. And the table's own height has nothing to do with this number — you drag its bottom edge to resize it.

3.11 Warm-start from library default on

The master switch for refine rounds. On, a non-explore round begins with your top-list champions dropped into generation 0 unchanged and the rest of the population filled randomly around them. Off, every round is a from-scratch exploration and Explore every N-th stops meaning anything. Leave it on for normal use — this is what turns a series of independent searches into something that compounds. Turn it off deliberately when you want breadth rather than a peak.

Watch out. A champion with more operations than the current Max. nodes is silently skipped when the round seeds itself — no message anywhere — so lowering Max. nodes after mining large formulas quietly converts your warm start back into a random start.

Part 04

The engine knobs — simulation, formula shape, breeding, fitness

How the simulator trades a formula, how big a formula may get, how the next batch is bred, and how candidates are scored. Every value here is read at Start and overrides config.ini. The boxes are spinners: only the arrows respect the limits — a typed number is passed through unchecked.

4.01 Target-vol (ann.) default 0.25 range 0.01–3.00

The annual volatility the simulated book is steered to. The engine keeps a running estimate of how much the book is swinging and scales every position so that swing lands on this number — this is the leverage knob, solved for rather than typed. Because the score is a Sharpe ratio, it barely changes WHICH formulas win: measured on one signal at 0.10 / 0.25 / 0.60 / 1.50, realised volatility came out 0.10 / 0.26 / 0.63 / 1.58 (the targeting works) while Sharpe stayed +0.65 / +0.69 / +0.67 / +0.63. What it changes is the risk you would actually run: worst drawdowns in the same test were –20% / –44% / –76% / –99.7%. Set it to the risk you are willing to hold, not as a tuning dial.

Watch out. There is no margin model and no liquidation, so the high end is arithmetic rather than a tradeable setting. Across those same four targets the CAGR of the identical formula went 6.4% → 15.6% → 25.1% → –21%: at 1.50 the Sharpe was still fine and the compounded result was a loss, purely from the size of the swings.

4.02 Fee (turnover) default 0.001 (10 bp) range 0.0000–0.0500

What trading costs, as a fraction of the notional that changes hands, charged each time the book moves — so a round trip costs about twice this. It is charged on the CHANGE in a position, not the position: a formula that sits still pays nothing. This is the only brake the search has on churn; there is no separate turnover penalty anywhere in the fitness. Set it at or above your real all-in cost, fee plus expected slippage, and the search will refuse to hand you signals that only work for free.

Watch out. Setting it to 0 does not merely flatter results, it changes what wins. Measured: a fast-turning signal scores +0.57 with fees off and –0.76 at the default, while a slow one only drops +1.13 → +0.69. A whole family of high-churn formulas becomes "profitable" at zero, and those are what the search then hands you. (A hard-coded 10% no-trade band already absorbs small rebalances, so your real cost is below turnover × fee.)

4.03 Max. depth default 6 range 2–12

How many operators may be stacked on top of one another before a branch must end in a raw feature. This binds harder than it looks: of the 25 operators, 19 take a single argument and only 6 take two, so formulas grow as long narrow chains rather than wide bushes. In a real library of 270 champions only about one node in ten was a two-argument operator and roughly a third of the champions sat exactly at the depth ceiling — depth is usually what caps a formula, not Max. nodes. Drop it to 3–4 for formulas you can read at a glance; raise it to 8 to allow longer chains of transformations.

Watch out. Depth is enforced by pruning, and pruning replaces an over-deep branch with a random raw feature — a child that breaks the limit is not rejected and retried, it is damaged and evaluated anyway. Warm-start champions are checked for size but never for depth, so lowering this after mining deep formulas means every child bred from them gets random leaves grafted in.

4.04 Max. nodes default 22 range 3–80

The total count of everything in the formula, operators and raw features alike; `ts_mean:30(close)` is 2 nodes. At the 22-node default a formula is typically 16–17 operations over 3–5 features. This is the size limiter and the number the complexity penalty is charged against. Small (10–15) gives simple, readable, harder-to-overfit formulas and quick convergence; large (40–80) gives expressive combinations and many more ways to memorise TRAIN. If you raise it, raise Complexity penalty too, or you are paying formulas to be complicated.

Watch out. The limit leaks exactly where it matters: crossover and mutation children over the limit are rejected and retried, but the fresh random injections are added with no size check at all. Measured across 20 seeds, at the default this is negligible — at Max. nodes 8, roughly 4 of the 10 injections per generation are oversize, and they can be selected as parents.

4.05 Tournament size default 5 range 2–50

How a parent is chosen: draw this many candidates at random and keep the best. The single dial for selection pressure. Low (2–3) makes mating nearly random — diverse, exploratory, slow to improve. High (15–50 against 200) lets the top few per cent dominate at once: fitness climbs for two or three generations, then the population becomes copies of one lineage. Watch the per-generation "evaluated N (uniq M)" line: if the unique count stops growing, pressure is too high for your population.

Watch out. Draws are made with replacement, so at 50 out of 200 the current best is in nearly every tournament and selection has become "always breed the best". The fitness compared here is the penalised one — a formula can lose for being large or for resembling a champion, rather than for trading worse.

4.06 Elitism default 6 range 0–50

How many of the best formulas are copied into the next generation completely untouched. Without it, the best thing you have found can be destroyed by the very mutation meant to improve it. 6 out of 200 (3%) is sane. At 0, progress can go backwards; above roughly 10% of the population the same handful keeps re-winning tournaments and the search stalls. Elites cost nothing to evaluate, but every elite slot is one not spent breeding something new.

Watch out. Elites are picked by the penalised fitness of that moment, similarity against the Hall of Fame included — so a formula can be an elite in one generation and drop out in the next without changing at all, simply because a champion resembling it entered the Hall of Fame in between.

4.07 Random/generation default 10 range 0–200

How many brand-new random formulas are dropped into each generation. Crossover and mutation can only recombine what the population already holds; this is the only channel that reintroduces operators and features the gene pool has lost. 10 out of 200 keeps a trickle of novelty. Raise it (20–40) when the

unique-formula count flattens early. Every injection is a full evaluation and nearly all are junk, so it is a direct CPU tax.

Watch out. Elitism + Random/generation is not capped by Population: 50 and 200 against a population of 200 silently produces 250 individuals, costing 25% more CPU than you asked for. And when those two together already reach Population, no crossover or mutation happens at all — the run degenerates into pure random search.

4.08 Crossover share default 0.60 range 0.00–1.00

Of the children bred each generation, the fraction made by crossover — cutting a branch out of one parent and grafting it into a copy of another. The rest are mutations of a single parent, in one of three equally likely kinds: replace a branch with a fresh random one; swap an operator for another of the same kind keeping its horizon (ts_mean:30 → ema:30); or nudge a lookback window, usually a small jitter and about one time in five a jump to the coarse starting grid. Crossover recombines whole working ideas; mutation does the fine work and is the only thing that tunes lookback windows.

Watch out. At 1.00 window tuning is off entirely — the nudge lives only inside mutation. The only fine tuning left is the polish pass after evolution ends, which touches just the top 5 champions.

4.09 Complexity penalty default 0.010 per node range 0.000–1.000

Subtracted in proportion to node count: score – penalty × nodes. At the default the whole spread from the smallest formula to a full 22-node one is about 0.2 Sharpe — meaningful, not decisive. At 0.05 a full-size formula gives up 1.1 Sharpe and the search collapses toward stumps. Raise it when champions come out unreadable and their TEST lags their fitness (the classic overfit signature); lower it toward 0.005 when you have raised Max. nodes and want the room used.

Watch out. This steers selection, not the output: the Hall of Fame is sorted and deduplicated on the RAW score, so a bloated formula with the best raw score still tops your library however high you set this. Under "Optimize by win rate" it is multiplied by 0.1 internally, so the same typed number is ten times gentler.

4.10 Correlation threshold default 0.70 range 0.00–1.00

How alike two formulas' results must be before the app calls them the same alpha — the correlation of their per-bar TRAIN+VAL returns, ignoring sign. It does two jobs: it triggers the similarity penalty during selection, and it is the deduplication rule for the Hall of Fame. Lower it (0.5–0.6) when your library fills with near-identical variants of one idea; raise it (0.8–0.9) when too few distinct families come back and you would rather keep close cousins than nothing. Because sign is ignored, a formula and its exact inverse count as duplicates.

Watch out. Both ends are broken, in opposite ways. At 0.00 every candidate counts as a duplicate of its nearest champion — fed 500 unrelated candidates the list ended with exactly one entry (measured). At 1.00 dedup can never fire and the Hall of Fame degenerates into fifteen twins of the same formula.

4.11 Similarity penalty default 0.5 range 0.0–2.0

How hard a candidate is punished for resembling a champion already in the Hall of Fame: $\text{penalty} \times (\text{correlation} - \text{threshold})$, against the most-similar champion. It exists to stop the population converging on one idea. Raise it when every round returns variations of the same alpha; lower it to 0 when you are deliberately refining one family. Costs no CPU.

Watch out. At the defaults this is close to cosmetic: the largest subtraction possible is $0.5 \times (1 - 0.70) = 0.15$ Sharpe, noise next to a champion Sharpe of 1.5–2.5. To make diversity actually cost something, lower Correlation threshold as well. Setting it to 0 does not switch deduplication off — that is a separate mechanism on the threshold.

4.12 Hall of Fame size default 15 range 1–100

How many champions a round produces — the formulas written into your library when it finishes. Not a simple top-N: a new formula correlating above the threshold with an existing champion replaces it if it scores better and is discarded if it does not, so the list stays diverse rather than merely high-scoring. Raise it when rounds keep reporting "none kept" and you want the near-misses; lower it (5–8) when the library fills with entries you never open. Every formula is correlated against every champion twice per generation, so $15 \rightarrow 100$ makes that bookkeeping about seven times heavier.

Watch out. Raising it does not get you more of your best work: the polish pass that fine-tunes champion lookbacks is hard-coded to the top 5 whatever this says, and champions already in your library are skipped when the round writes out — so a larger Hall of Fame mostly adds weaker, newer entries.

4.13 Robust blocks (0 = legacy) default 0 range 0–12

An alternative way of scoring. At 0, fitness is $\min(\text{TRAIN}, \text{VAL})$ Sharpe. At 2 or more, the app merges TRAIN and VAL into one span, cuts it into that many equal time slices, shrinks each slice's Sharpe by its own margin of error, and scores the formula on the near-worst slice. The intent is that a formula must work across regimes. The honest answer: the app's own config.ini records a retrospective A/B on an 861-alpha daily library that found no improvement — rank correlation with held-out TEST was +0.40 legacy, +0.40 for the best 3-block variant, +0.26 for 5 blocks. $\min(\text{TRAIN}, \text{VAL})$ already IS a two-block worst-regime rule, split at the bull/bear boundary. Treat it as an opt-in experiment; if you try it, use 3.

Watch out. Turning it on changes what "fitness" means everywhere and nothing in the interface says so. Block fitness routinely runs deep in the negative (the same three formulas: $-0.28 / -0.32 / -0.89$ legacy versus $-0.86 / -1.92 / -1.26$ at 3 blocks), the leaderboard sorts on that raw number without knowing which rule produced it, so block-scored champions land under every legacy row and never seed a refine round. If you upgraded from a version predating this setting, the box opens showing 5 rather than 0 — check it before your first Start.

Part 05

Date segments, and the honesty discipline

Four dates define the three stretches. The search may stare at TRAIN and VAL as long as it likes; TEST is walled off and only ever reported — never used to pick, rank, seed or fine-tune. That single rule is the difference between a number you found and a number you fitted.

5.01 TRAIN start default **2019-09-05** range **YYYY-MM-DD · checked as you type and again at Start**

The very first bar the app loads, for any purpose. Nothing before it exists as far as AlphaNode is concerned. Earlier means more history and more market moods to prove a formula across, at the price of speed: every candidate is simulated over the whole TRAIN-start-to-TEST-end span on every evaluation (TEST included — simulated, just never scored), so doubling the span roughly doubles the cost of a round. There is a warm-up tax at this edge: a formula may look back 200 bars, rolling operators start halfway through their window, the first five bars are never tradable, and the volatility estimate needs 30 bars (2,880 at 15m) before it means anything. Leave a couple of hundred bars of runway before the part of TRAIN you care about.

This field decides the length of the whole window, so it is the one the per-timeframe ceiling argues with. TRAIN start to TEST end may not exceed 3,000 bars on 1d, 16,000 on 4h, 40,000 on 1h or 96,000 on 15m — roughly 8.2, 7.3, 4.6 and 2.7 years. The limits sit about 15% above each bar size's own recommended window, which is the span this build is tuned for; the cost of a round is linear in bars, and the same four calendar years is 1,460 daily bars but 140,000 fifteen-minute ones. Exceed it and a red line appears under the fields saying how many days to move TRAIN start, and Start refuses until you do.

Watch out. The download never asks for your TRAIN start — it fetches from the timeframe's own recommended date. Set TRAIN start earlier and nothing warns you: the missing stretch is filled with a frozen price, every formula sits flat through it, and because flat bars stay in the scoring window every Sharpe is dragged down. If under 10% of TRAIN is left trading, the activity filter rejects every candidate and the round finds nothing.

5.02 VAL start default **2021-11-01** range **strictly between TRAIN start and TEST start – enforced**

The wall between TRAIN and VAL, and the most powerful anti-overfitting lever in the app — a date, not a number. An alpha's score is the LOWER of its TRAIN and VAL Sharpe, so a formula that made all its money in one kind of market is judged by the kind where it did not. Put it where the market's character genuinely changes: the shipped 1d cut lands it at the top of the 2021 bull run, so TRAIN is the run-up and VAL the bear market that followed. Move it later and TRAIN gets richer while VAL thins into noise; earlier and VAL quietly becomes the dominant judge. Keep VAL long enough to mean something — the engine's own error bar puts about ± 1.2 around a one-year VAL measured at Sharpe 1.0, ± 0.87 at two years, ± 0.61 at four.

Watch out. Setting "Robust blocks" to 2 or more silently retires this date from the score: the engine re-cuts the whole TRAIN-to-TEST span into equal slices and the VAL boundary plays no part in the number any more. It still decides validity and the VAL row still shows on the champion card, so nothing looks different while the number you are optimising has changed meaning.

5.03 TEST start

`default 2023-01-01``range strictly between VAL start and TEST end – enforced`

The wall. Everything from here to TEST end is held out: it never enters the score, never ranks the shortlist, never seeds a warm start, never guides window fine-tuning. Treat its length as a budget you are spending — push it later and the search gets more data while your verdict gets noisier; pull it earlier and you buy a long, trustworthy TEST at the cost of a thinner search. The shipped 1d split is deliberately lopsided in TEST's favour: roughly 790 TRAIN, 430 VAL and 1,280 TEST bars. For a daily search keep TEST to at least 1.5–2 years of calendar.

Watch out. A too-short TEST degrades quietly, not loudly: under 5 traded bars the TEST cell just shows "—"; the direction columns each need 30 bars in their bucket and blank out first; and sorting by TEST OOS drops every row with no TEST number, so alphas appear to vanish. Re-cutting this date also leaves the library in mixed vintage — fitness and TEST OOS are frozen at mining time while the direction and activity columns are recomputed against today's dates, and nothing on screen says so.

5.04 TEST end

`default 2026-07-05``range strictly after TEST start – enforced; still not clamped to the data you actually hold`

The last bar the app loads, for everything — search, leaderboard numbers, portfolio build, CSV export. Normally this should be today, or the last complete bar in your snapshot; the fetcher defaults its own end date to today, so your data is almost always ahead of this field.

Watch out. Two traps. Setting it past the end of your data costs Sharpe silently: the last close is carried forward, those bars drop out of eligibility, every strategy returns exactly zero across them, and TEST Sharpe is multiplied by roughly the square root of the share of real bars. And the shipped default is a fixed date, not "today" — leave it alone for a few months and the most recent stretch of real downloaded history is simply never looked at.

5.05 The recommended cut at each bar size

1d — 2019-09-05 / 2021-11-01 / 2023-01-01 / 2026-07-05, about 790 / 430 / 1,280 bars over 6.8 years.
 4h — 2020-06-01 / 2023-01-01 / 2024-09-01 / 2026-07-01, about 5,660 / 3,650 / 4,010 bars. 1h — 2022-06-01 / 2024-06-01 / 2025-06-01 / 2026-07-01, about 17,500 / 8,760 / 9,480 bars. 15m — 2024-01-01 / 2025-04-01 / 2025-12-01 / 2026-07-01, about 43,800 / 23,400 / 20,400 bars. The pattern is deliberate: finer bars trade calendar span for bar count. Each bar size also carries a hard ceiling on the whole window — 3,000 / 16,000 / 40,000 / 96,000 bars — so the recommended cut uses 83% to 91% of what its timeframe allows. Switching the timeframe selector refills all four dates from this table, and a settings file whose dates are still some other timeframe's recommendation is refilled on launch (a window you typed yourself is never touched).

Watch out. The extra bars do NOT buy a more believable verdict — the error bar on a Sharpe depends on calendar YEARS, not on how finely you sliced them. At Sharpe 1.0 the shipped TEST windows carry roughly ± 0.65 (1d), ± 0.91 (4h), ± 1.18 (1h) and ± 1.60 (15m). A 15m TEST holds 20,000 bars and still cannot tell a Sharpe of 1 from a Sharpe of 0. Read intraday TEST as a sanity check, not as evidence, and lean much harder on the forward track there.

5.06 Working with a re-cut

Once you move TEST start or TEST end, the honest options are to wipe and re-mine on the new cut, or to treat the surviving library as historical and never compare its numbers against new rows. In practice: mine under one cut, read TEST once, enrol whatever you liked into the forward track — an enrolled entry freezes its own dates and settings, so it keeps accumulating genuinely unseen bars whatever you do to the fields afterwards — then push the segments forward and clear. Save a named session first: it snapshots the settings file too, so both the old library and the old cut can come back.

Watch out. A library accumulates rows mined under different settings, and a row records which objective it was mined under but not the dates. Two rows both reading "+1.42" may have been measured on entirely different years.

5.07 by — how portfolio members are picked

default TEST • not stored in settings, so every launch opens on TEST

range TEST | fitness | combo

Combining alphas is a selection step like any other, and this is where the discipline is easiest to break by accident. "fitness" picks members by the worse of TRAIN and VAL, so TEST never enters the choice and the combined TEST curve is genuinely out-of-sample — you can quote it. "combo" is better still: instead of the N best individuals it searches for the best COMBINATION of N, taking a pool of the top alphas by fitness ($4 \times N$, never under 12 or over 30), simulating the pool once, then running a greedy-then-swap search that maximises the mix's Sharpe on TRAIN+VAL only. Because the goal is the mix's Sharpe rather than any member's, it naturally hunts members that do not move together.

Watch out. The selector opens on TEST — the one option of the three that breaks the discipline. What is still real in a TEST-picked portfolio is the diversification GAIN: combined beating every member is a structural fact about blending uncorrelated signals. The LEVEL of the curve is not.

Part 06

The leaderboard, column by column

Four columns — fitness, TEST OOS, maxDD, CAGR — are stored numbers written when the formula was found. The rest are recomputed on demand by a background helper that re-simulates the formula on TEST only, using your CURRENT pairs, target vol, fee and dates. A dot (·) means "still computing"; an em dash (—) means "no honest number here".

6.01 ★ default always shown

A star means the formula is in your Favorites. Clicking anywhere in the cell toggles it — no dialog. Use it as your shortlist while scrolling a library of hundreds.

Watch out. The star set is keyed by the 6-char ID with no timeframe in the key, so the same formula text starred on 1d also shows starred on 4h. Sealed rows cannot be starred at all.

6.02 # default always shown

The row's position in the current sort — 1 is the top row of this sort, not a permanent rank. A place-marker for talking about rows, never a quality score, and not sortable.

6.03 fitness default always shown · the default sort

The score the search ranked by: the WORSE of the TRAIN and VAL Sharpe. The only column that is honest for picking — TEST never touched it. As a rough rule on daily bars, above roughly +1.5 is worth opening; well above +2.5 is usually too good and worth suspecting rather than celebrating.

Watch out. The tooltip does not mention that a concentration guard is on by default and docks this number when the book averages under 3 meaningful positions, nor that the value shown is BEFORE the size and similarity penalties the engine uses to pick tournament winners. With "Optimize by win rate" on the cell shows a percentage, and clicking this header sorts rows mined under your ACTIVE objective to the top as a block — a 0.57 win rate and a +1.8 Sharpe are not on the same ladder.

6.04 TEST OOS default always shown

The Sharpe on the held-out slice — the honest out-of-sample verdict on a single formula. A dash means too little to measure: fewer than 5 TEST bars where the strategy actually made or lost anything. Read it, do not select on it. Expect it to sit well below fitness; a formula whose TEST is CLOSE to its fitness is the rare good sign, and a wide gap is the classic overfitting tell.

Watch out. Clicking this header does far more than re-sort: the whole library is re-read and re-ranked by held-out TEST, every row with no TEST number is dropped, and the heading changes to "cherry-picked △". That is peeking by construction. Click "fitness" to get the honest population back.

6.05 maxDD default shown by default

Worst peak-to-trough fall of the TEST equity curve, as a negative percentage. The pain metric — this is what decides whether you could actually hold the thing. Two formulas with the same Sharpe and drawdowns of −15% versus −45% are completely different products. It moves with your target-vol setting: halve the target and the drawdown comes down roughly in step.

Watch out. Unlike the recomputed cells, this prefers the value STORED when the alpha was mined. Change target vol, fee, pairs or dates and maxDD and CAGR still describe the OLD settings while every other recomputed cell in the same row describes the new ones.

6.06 CAGR default shown by default

Compound annual growth rate on TEST. A dash means the equity ended at or below zero, so no growth rate exists. Pair it with maxDD, never read it alone: $\text{CAGR} \div |\text{maxDD}|$ is a quick sense of return per unit of pain. Like maxDD it moves with your target-vol dial, so it says as much about your risk setting as about the alpha's edge.

Watch out. Same stale-value trap as maxDD — taken from the library row first.

6.07 sortino default hidden – right-click a header to switch it on

Like Sharpe, but only losing bars count toward the wobble in the denominator; upside volatility is free. Look at it when a Sharpe seems unfairly low: a big sortino-to-Sharpe gap means the volatility Sharpe punished was mostly good volatility. If it is barely above Sharpe, the returns are symmetric and Sharpe was telling the truth.

Watch out. A dash here can mean the opposite of failure — with no losing bars at all the downside deviation is zero and the code returns nothing rather than infinity. In practice that almost always means an alpha that traded a handful of bars; cross-check tr/yr·a.

6.08 T ↑ default shown by default · the header carries the bucket's bar count

TEST Sharpe measured only on bars where the market as a whole was trending UP. The regime is decided by the market, not the formula: an equal-weight basket of your universe has its drift measured over the trailing ~30 calendar days, and if that drift is statistically confident ($t \geq 1.28$) the bar takes the drift's sign. This is the regime question — does the alpha only make money when everything rises, in which case you have bought beta, not alpha? T ↑ and T ↓ both meaningfully positive is genuinely direction-neutral and much more valuable.

Watch out. The label is deliberately lagged one bar; without that lag a plain market-clone showed roughly +1.3 of pure T ↑-minus-T ↓ Sharpe out of thin air. Two different causes of a dash: under 30 TEST bars in the bucket (hits every row at once, since the split is the market's calendar) or the formula sat flat through the whole bucket (per-row).

6.09 T ↓ default shown by default

The same measurement on confidently DOWN-trending bars. The most informative of the three for crypto: most formulas quietly make their money long, and a positive T ↓ is the evidence that the alpha survives a bear leg. If you plan to run this through a drawdown, read this column first.

6.10 T ~ default shown by default

TEST Sharpe on FLAT bars — where the basket's drift was not confident either way. Usually the largest of the three buckets. This is where mean-reversion and relative-value formulas earn their keep and where trend-followers bleed. Strong in T ~ and weak in T ↑/T ↓ is a chop specialist: useful, but it underperforms exactly when the market is most exciting.

6.11 L/S /yr·a default hidden – right-click a header to switch it on

Long entries and short entries, per asset, per year, on TEST — "3.2/1.8" is roughly three new longs and two new shorts per coin per year. An entry is the moment the target weight crosses into long or short from flat or the opposite side. The alpha's directional personality: 4.0/0.1 is a long-only formula wearing a long/short costume and will die in a bear market whatever its TEST Sharpe says; near 2.5/2.5 is genuinely two-sided. Normalised per asset, so a 5-coin universe compares directly with a 30-coin one.

Watch out. Flipping straight from long to short counts as ONE entry, and holding for eight months counts as zero further entries — this measures decision frequency, not how much notional you push around.

6.12 tr/yr·a default hidden – right-click a header to switch it on

Trades per asset per year on TEST: the two L/S numbers added together. The activity dial. Low single digits means a slow position-holder where fees barely matter; a few hundred means an alpha whose profitability lives or dies on your fee setting — re-check TEST Sharpe with a realistic fee before believing it.

Watch out. It counts INTENDED entries, not fills: the simulator leaves a position alone whenever the required change is under 10% of the new target. Honest about how often the formula changes its mind, an overstatement of how often you would touch the exchange.

6.13 win% default shown by default

The share of bars on TEST where the strategy made money. The denominator counts only ACTIVE bars — flat bars are excluded entirely. Best read against tr/yr·a: a high win% on an alpha that trades constantly is a real edge, a high win% on one that trades rarely is a small sample dressed up as skill. Win rate and profitability are different questions — 45% with big winners beats 65% with big losers.

Watch out. The tooltip says "profitable days" but on 4h/1h/15m these are bars. And do not read a win-rate-mined row's fitness cell as a corrected version of this one: fitness is measured on TRAIN and VAL, this column on TEST. Different periods, not two views of the same number.

6.14 win ↑ default shown by default

Win rate on confidently up-trending bars — same regime and same one-bar lag as T ↑. Use it with T ↑ to separate two failure modes: decent win ↑ with poor T ↑ means the formula wins often but its losses are large; poor win ↑ with decent T ↑ means it wins rarely and wins big. Both are real strategies, needing very different sizing and very different patience from you.

Watch out. A stricter evidence bar than the T columns: 30 bars in the regime AND at least 5 where the formula made or lost something. A dash here beside a numeric T ↑ just means it was mostly flat during up-trends.

6.15 win ↓ default shown by default

Win rate on confidently down-trending bars, same definition and the same two evidence floors. The stress-test companion to T ↓: if win ↓ collapses while win ↑ stays high, the formula is a long-bias machine whatever its entry counts suggest.

6.16 ID default shown by default

A stable 6-character fingerprint of the formula (the first six of its md5), sitting immediately after # at the left edge of the table. The same formula always gets the same ID. This is the handle used everywhere else — the forward track, favourites, the search box, both CSV exports.

Watch out. The two exports do not print it the same way: "Export table" writes the 6 characters you see, "Export full library" writes 12, so a plain text match between the files fails unless you trim.

6.17 formula default always shown

The alpha itself, in full. The column is sized to the widest formula rather than truncated, so long ones bring out the horizontal scrollbar. Length is itself information: the search subtracts a penalty per node, so a short formula that still made the library beat that handicap and is usually the more robust one. Click a cell to make the text selectable, or use Copy formula.

Watch out. Absolute SCALE is meaningless — a formula and the same formula times 100 are literally the same strategy, so do not read big constants as "aggressive". A locked row shows "locked" here and dashes in all nine recomputed cells, while fitness, TEST OOS, maxDD and CAGR still show real numbers, because those were stored before it was sealed.

Part 07

What you can do with a formula

Everything reachable from a leaderboard row: the equity window and the five buttons inside it, the table's own menus and exports, and the two buttons that can destroy work.

7.01 Double-click a row → the equity window

Growth of \$1 on a log scale, shaded into TRAIN | VAL | TEST, with an equal-weight buy-and-hold basket of your universe alongside; all curves net of fees. Stretches where market volatility was above its own one-year median are tinted, and a lower panel tracks the unrealised profit of open positions. The header repeats Sharpe / CAGR / drawdown per segment. You are looking for a curve that keeps its SHAPE across the TRAIN→VAL→TEST boundaries rather than one that made its money in a single lucky stretch. It is a live matplotlib canvas — wheel to zoom, drag to pan.

Watch out. A double-click landing inside the ★ cell never opens the chart: each of its two clicks toggles the star, so nothing appears to happen. If a double-click does nothing, you hit the first column.

7.02 Passport

Builds an explanation of the formula rather than a measurement of it: the expression as a labelled tree, plain-English reading steps, the signal plotted against one asset's price, an ablation showing which inputs it really depends on, its holding period and long/short balance, and which classic archetype its returns most resemble. No network, no AI — a fixed translation of the operators. Use it when a formula looks good and you want to know WHY before trusting it: an alpha you can describe in a sentence is one you can reason about when it stops working.

Watch out. An explanation is not evidence. A convincing story does not make the TEST number more real — the ordering is always TEST first, forward track second, story third.

7.03 PDF report

A four-page analytics dashboard written where you choose: headline numbers with equity and drawdown; exposure, long/short balance and turnover over time; the structure of the weights plus a monthly-return calendar; and a TRAIN/VAL/TEST breakdown with attribution and written conclusions. Every page is stamped with the date, the time and the 6-char ID. This is the artefact to keep or hand to someone else — the app's state can be cleared, a dated PDF cannot. The Portfolio card has its own button that does the same for a combined book.

7.04 Serve signal (API)

Starts a small local web service that keeps this formula's live target positions available as JSON: it fetches fresh closed candles from Binance for your universe, recomputes the weights and serves the latest bar at /signal, with /health reporting its age. It claims the next free port from 8799 and appears in the SIGNAL API card. This is the bridge to an execution bot — point your code at `http://127.0.0.1:<port>/signal`. It uses the same target vol and fee as the leaderboard, so what it serves matches what you were looking at; it refreshes every 15 minutes on 1d and 4h, every 5 on 1h and 15m.

Watch out. On daily bars the response carries "leverage" separately and the weights are the relative book — multiply to get real sizes. On intraday the weight already includes leverage and the served leverage is a flat 1.0. It is weight × leverage × equity that stays identical between the two, not the weights.

7.05 Download signals (CSV)

The full position history: one row per asset per bar where it held a meaningful position, labelled with its segment, the ticker, LONG or SHORT, the signed weight as a number and a percentage, and that asset's OHLCV on the bar. A dialog then shows what the strategy would be holding on the very last bar — the quick "what do I buy right now" answer. This is how you take the strategy out of AlphaNode.

Watch out. These weights are the normalised target book (absolute values sum to 1 across a bar) taken BEFORE volatility targeting. They are relative allocations, not position sizes — size directly from this column and you are running at whatever leverage you happened to pick.

7.06 Forward track +

Enrols this alpha: the strategy is frozen exactly as it stands — formula, universe, target vol, fee, timeframe — and from then on paper-traded once per closed bar on freshly downloaded Binance data from a clean \$10,000. Do this the moment a formula survives TEST and looks explainable. Everything else in the app is measured on data that already existed; this is the only clock that runs forward. Give it weeks before reading anything into the number.

Watch out. Enrolling the same formula twice is refused rather than starting a second track. Delete removes the strategy AND its whole paper history permanently. On 4h/1h/15m, steps only happen while the app is open and fees weigh proportionally more on short bars — the enrol dialog warns you.

7.07 Chart (Forward Track panel)

The forward equity curve of an enrolled strategy: one point per live paper step since enrolment, with the starting capital as a dashed line. This is the only curve in the app that was never fitted to anything — every point was computed on data that did not exist when the formula was found. It is short and noisy for weeks, and that is fine; the value is that it cannot lie to you.

Watch out. History is append-only and nothing is recomputed backwards: if the app was closed for a week, the gap stays a gap in the line.

7.08 ★ Favorites

The button in the leaderboard heading opens your starred list: date added, ID, the timeframe it came from, fitness, TEST OOS and the formula. Double-click for the equity chart. The starred formulas live in their own file outside the library, so they survive "Clear all history": the whole row is copied at star time, which is what keeps it readable after the library row is gone. They belong to the SESSION, though — saving a session saves its stars, and loading one replaces the stars on screen with that session's. A star points at a formula in a particular library; inheriting it into a workspace mined on another basket, another cut or another timeframe said nothing true. Star anything in this session you would be upset to lose, and save the session to keep it.

Watch out. Loading a session archive saved by an older build leaves you with no stars at all — those archives carry no favorites file, and a load replaces every session-owned file wholesale rather than merging. A favourite is also a snapshot: the whole row is copied at the moment you star it and never refreshed. Change target vol, fees or dates and the leaderboard's numbers move while the Favorites window's do not.

7.09 Copy formula / Copy formula + metrics

Right-click a row for both. The plain copy (also Ctrl+C) puts the bare expression on the clipboard — use it when pasting into code. The +metrics version adds a line with the fitness value, a note if it was mined by win rate, and the TRAIN, VAL and TEST Sharpe ratios — use it when recording a candidate in notes, since that one line stays readable months later.

7.10 Right-click a header → the Columns menu

default on: ID, maxDD, CAGR, T↑, T↓, T~, win%, win↑, win↓ · off: sortino, L/S, tr/yr

A checklist of the twelve optional columns; your choice is saved and reused next launch. ★, #, fitness, TEST OOS and formula are always present. Turn L/S and tr/yr on the moment you start caring whether an alpha is genuinely two-sided or how much its fee assumption matters.

Watch out. Column ORDER is fixed and follows the menu, not the order you tick things in — so switching sortino on inserts it mid-table rather than appending it at the right. ID is the one exception: it is pinned immediately after #, because it names the row instead of scoring it. Hiding a column only hides it: sorting and both CSV exports carry every column regardless.

7.11 Sorting default opens on fitness, descending – the sort is not remembered between launches

Clicking a header sorts by it, best first; clicking again flips. Everything except ★, # and ID is sortable, formula included (alphabetically). Only fitness and TEST OOS change WHICH formulas the table contains; every other column simply reorders what is already there.

Watch out. Sorting by any recomputed column forces the app to compute that statistic for EVERY row in the table, not just the visible ones — on a large library that is a long job with a stalled-looking table. And rows with no value are treated as minus-infinity, so ascending sort puts all the dashes at the TOP.

7.12 Export table (CSV)...

Writes exactly the rows currently in the table, in the current order, with twenty-two columns — every leaderboard metric plus the TRAIN and VAL Sharpe and the objective the row was mined under — regardless of what is hidden on screen. The export for a specific working set: filter with the search box, sort how you want, then export those rows with all their analysis columns.

Watch out. The recomputed columns come from an in-memory cache holding only rows that have actually been on screen, filled roughly a screenful at a time. Export a 600-row table without scrolling and everything past the first screenful comes out blank in those cells. Scroll first, or sort by one of them to force the whole set.

7.13 Export full library (CSV)...

The CSV button in the heading: every alpha ever mined into this timeframe's library — no family dedup, no TEST filter, no on-screen subset — sorted by fitness, in nineteen columns including the round it was found in, its timestamp and the stored Sharpe / drawdown / CAGR / active-bar count for each of TRAIN, VAL and TEST. The archive export. It reads the library file directly, so it works while the node is mining.

Watch out. It carries NONE of the recomputed columns — no sortino, no regime splits, no trade counts, no win rates. And it sorts strictly by the raw fitness number, so in a mixed library every win-rate row (a value under 1) sinks below every Sharpe row regardless of quality.

7.14 ► Build portfolio default top 6, by TEST · neither is saved – both reset every launch

Runs N alphas together as one combined book and shows the combined curve against buy-and-hold. Once a build finishes the member list fills in under the metrics — which alphas went in and how each does alone — and four buttons unlock: CSV (combined signals), Serve (one API for the whole book), PDF, and Track (enrol the whole portfolio into the forward track). Prefer "fitness" or "combo" over the default — see the entry on the by selector.

Watch out. Because neither setting is persisted, a habit of building by fitness has to be re-established every session. The member list joins fitness and TEST OOS from the library by formula text, so after "Clear all history" a still-built portfolio keeps listing its members with dashes in those two columns — the portfolio outlives the library it came from, and double-clicking such a member cannot chart it.

7.15 "families only" switch default off – every alpha shown

Collapses the table to the best representative per family, where two formulas are the same family if their text is 80% or more similar. Turn it on when the search has converged and the top of your table is fifty cosmetic rewrites of one expression; turn it off to see the whole population or to search by ID. Both sets are cached in the background, so the switch is instant — and unlike the sort order, this choice is remembered across restarts.

Watch out. Families mode is capped twice: it stops at 20 distinct families and only ever examines the top 500 rows while collecting them. A genuinely novel formula at rank 600 is invisible here — it is a view of the head of the list, not a summary of the library.

7.16 Clear all history

The red button in the settings panel. Deletes the mined alpha libraries, the round history, the status file, the built portfolio and the cached equity images, after showing a count of what will go. Your search settings and your starred favourites survive. It refuses to run while the node is mining.

Watch out. It clears EVERY timeframe, not just the one you are looking at, while the counts in the dialog are read from the CURRENT timeframe only — so a dialog saying "40 found alphas" can be about to delete several thousand. This is the single most destructive button in the app and neither its label nor its dialog says so. Save a session first.